

DS-UA 111 HW4

August 17, 2023


1 Question 1

- a) False
- b) True
- c) True
- d) True
- e) False

2 Question 2

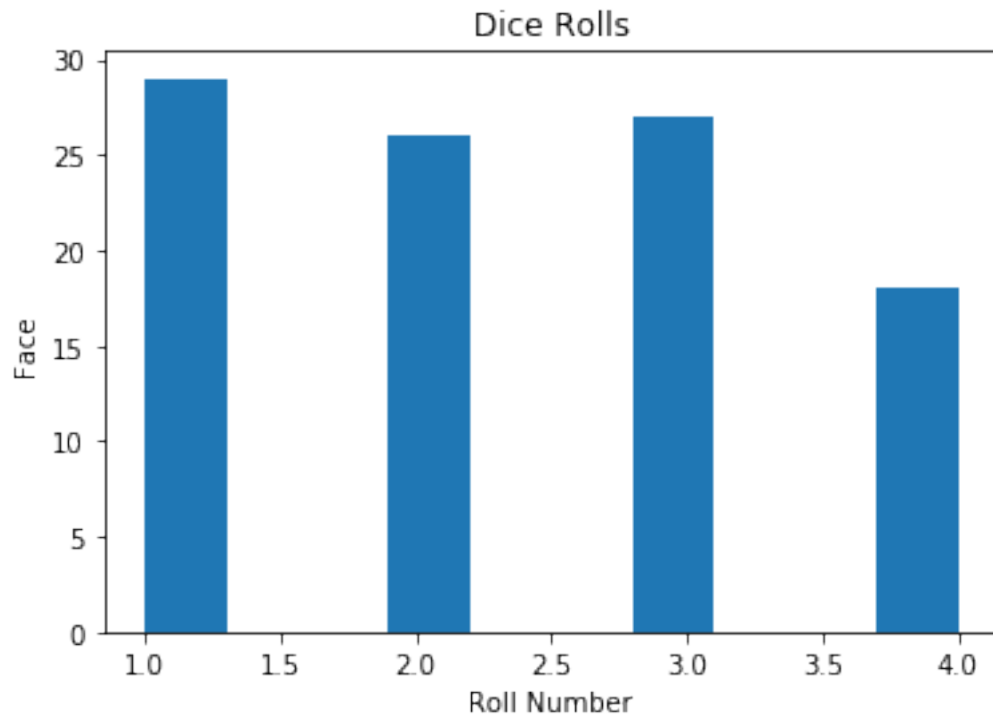
```
[1]: # part a
import numpy as np
import matplotlib.pyplot as plt

faces = np.random.choice(range(1,5), 10)
faces
plt.hist(faces)
plt.title("Dice Rolls")
plt.xlabel("Roll Number")
plt.ylabel("Face")
plt.show()
```

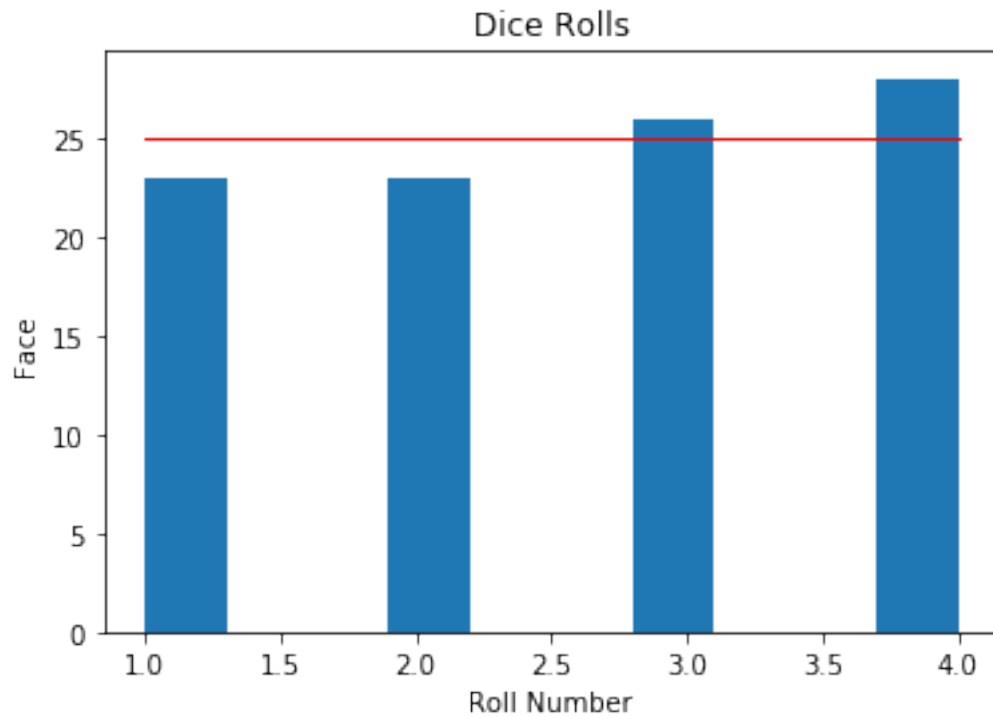


b) 0.2 ; 0.25


```
[2]: # part c
faces = np.random.choice(range(1,5), 100)
faces
plt.hist(faces)
plt.title("Dice Rolls")
plt.xlabel("Roll Number")
plt.ylabel("Face")
plt.show()
```



```
[3]: # part d
faces = np.random.choice(range(1,5), 100)
faces
plt.hist(faces)
plt.title("Dice Rolls")
plt.xlabel("Roll Number")
plt.ylabel("Face")
x_coord = [1,4]
y_coord = [25,25]
plt.plot(x_coord, y_coord, color="red", linewidth=1)
plt.show()
```



```
[4]: # part e
faces = np.random.choice(range(1,5), 10000)
faces
plt.hist(faces)
plt.title("Dice Rolls")
plt.xlabel("Roll Number")
plt.ylabel("Face")
x_coord = [1,4]
y_coord = [2500,2500]
plt.plot(x_coord, y_coord, color="red", linewidth=1)
plt.show()
```



- f) This happens because of the law of large numbers. The law of large numbers states that the higher the sample size of an experiment, the closer our empirical distribution matches the theoretical distribution.

3 Question 3

```
[5]: # part a
import pandas as pd
population = pd.read_csv("population_data.csv", squeeze = True)
pd.DataFrame(population).describe()
```

```
[5]:
```

	0
count	100000.000000
mean	8.330766
std	6.291950
min	0.000415
25%	3.314054
50%	7.047814
75%	11.993263
max	46.983332

```
[6]: # part b
population_sample = pd.DataFrame(population).sample(n=25)
```

```
population_sample.describe()
```

```
[6]:          0
count  25.000000
mean    8.542476
std     4.696871
min     0.424520
25%     4.647740
50%     8.547478
75%    11.374711
max    17.016525
```

```
[7]: # part c
population_sample = pd.DataFrame(population.sample(n=1000))
population_sample.describe()
```

```
[7]:          0
count  1000.000000
mean    8.354364
std     6.477645
min     0.003483
25%     3.345058
50%     6.906624
75%    11.602347
max    36.584645
```

d) Yes; this is because of the law of large number.

```
[8]: # part e
def percentile(input_data):
    y1 = np.power(input_data,0.25)
    y2 = np.sqrt(y1) * -1
    y3 = np.exp(y2)
    y4 = np.sin(y3)
    return np.percentile(y4, 80)
```

```
[9]: # part f
percentile(population)
```

```
[9]: 0.3177257855396434
```

```
[10]: # part g
population_sample_100 = pd.DataFrame(population.sample(n=100))
percentile(population_sample_100)
```

```
[10]: 0.3229746322037541
```

```
[11]: # part h
population_sample_1000 = pd.DataFrame(population.sample(n=1000))
percentile(population_sample_1000)
```

```
[11]: 0.31977613407151495
```

The complex statistic on a randomly chosen sample of size 1000 from the population is closer; LLN seems to work well when the statistics are complex.

4 Question 4

```
[12]: # part a
from scipy.stats import norm
housing = pd.read_csv('house.csv')
sale_price = housing['SalePrice']
pd.DataFrame(sale_price.describe())
```

```
[12]:          SalePrice
count    2930.000000
mean    180796.060068
std      79886.692357
min      12789.000000
25%     129500.000000
50%     160000.000000
75%     213500.000000
max      755000.000000
```

```
[13]: # part b
p = 0.01
crit_val_two_tail = norm(0,1).ppf(1 - p/2)
crit_val_one_tail = norm(0,1).ppf(1 - p)
print(crit_val_two_tail)
print(crit_val_one_tail)
```

```
2.5758293035489004
2.3263478740408408
```

```
[14]: # part c
h0 = 150000
z = (np.mean(sale_price) - h0) / (np.std(sale_price) / np.sqrt(len(sale_price)))
print(abs(crit_val_two_tail) < abs(z))
```

```
True
```

I reject the null hypothesis.

```
[15]: # part d
print(crit_val_one_tail < z)
```

True

I reject the null hypothesis.

```
[16]: # part e
h0 = 200000
z = (np.mean(sale_price) - h0) / (np.std(sale_price) / np.sqrt(len(sale_price)))
p_value = 1 - norm(0,1).cdf(z)
print(p_value)
```

1.0

I don't reject the null hypothesis.

```
[17]: # part f
p_value = norm(0,1).cdf(z)
print(p_value)
```

5.067385832657905e-39

I reject the null hypothesis.

```
[18]: # part g
p = 0.01
alpha = abs(norm(0,1).ppf(p/2))
a = np.mean(sale_price) - alpha * np.std(sale_price)/np.sqrt(len(sale_price))
b = np.mean(sale_price) + alpha * np.std(sale_price)/np.sqrt(len(sale_price))
a,b
```

```
[18]: (176995.18508695194, 184596.93504956685)
```

5 Question 5

```
[19]: # part a
import statsmodels.formula.api as smf
world = pd.read_csv("World.csv")
X = world[['gdp_10_thou', 'pop_total', 'spendhealth', 'fertility']]
y = world['lifeex_total']
life_expectancy_regression = smf.ols('lifeex_total ~ gdp_10_thou + pop_total +
    ↪ spendhealth + fertility',
    ↪ data = world).fit()
life_expectancy_regression.summary()
```



```
[19]: <class 'statsmodels.iolib.summary.Summary'>
```

```
"""
                                OLS Regression Results
=====
Dep. Variable:          lifeex_total    R-squared:                0.712
Model:                  OLS            Adj. R-squared:           0.704
Method:                 Least Squares   F-statistic:              88.53
Date:                  Thu, 17 Aug 2023 Prob (F-statistic):       1.06e-37
Time:                  09:50:01         Log-Likelihood:           -456.37
No. Observations:      148             AIC:                     922.7
Df Residuals:          143             BIC:                     937.7
Df Model:               4
Covariance Type:       nonrobust
=====
                        coef    std err          t      P>|t|      [0.025    0.975]
-----
Intercept             80.5665      1.666     48.366     0.000     77.274     83.859
gdp_10_thou           2.7797      0.598      4.652     0.000      1.599      3.961
pop_total            -0.0005      0.003     -0.171     0.864     -0.006      0.005
spendhealth           0.0375      0.310      0.121     0.904     -0.575      0.650
fertility             -4.7158      0.346    -13.616     0.000     -5.400     -4.031
=====
Omnibus:                 38.137    Durbin-Watson:           2.158
Prob(Omnibus):            0.000    Jarque-Bera (JB):        76.368
Skew:                    -1.144    Prob(JB):                2.61e-17
Kurtosis:                 5.674    Cond. No.                 613.
=====
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
"""
```

```
[20]: # part b
health_spending_increase_coeff = 1000000
lifeex_increase = health_spending_increase_coeff * 0.0375
print(lifeex_increase)
```

37500.0

c) 71.2%

d) 80.5665

e) Yes, we can reject the null hypothesis. The p-value for GDP is 0.000, which is less than our p-value threshold of 0.05.


```
[21]: # part f
p = 0.05
```

```
alpha = abs(norm(0,1).ppf(p/2))
a = np.mean(world['fertility']) - alpha * np.std(world['fertility'])/np.
↳sqrt(len(world['fertility']))
b = np.mean(world['fertility']) + alpha * np.std(world['fertility'])/np.
↳sqrt(len(world['fertility']))
print(a,"and",b)
```

2.5367014919366944 and 2.9772009470876957

g) Yes, because the model assumes a linear relationship.

```
[22]: # part h
from sklearn import linear_model
pred_y = life_expectancy_regression.predict(X)
residual = world['lifeex_total'] - pred_y
plt.scatter(pred_y, residual)
plt.show()
```



There doesn't seem to be a clear trend.

```
[23]: # part i
df = {"prediction": pred_y, "residual": residual}
pred_residual_corr = pd.DataFrame(df)
print(pred_residual_corr.corr())
```

```
           prediction      residual
prediction  1.000000e+00 -1.590993e-15
```

```
residual    -1.590993e-15   1.000000e+00
```

It suggests the presence of potential problems with the models doesn't depend on the prediction of life expectancy.

```
[24]: # part j
df = {"gdp_10_thou": world['gdp_10_thou'], "pop_total": world['pop_total'],
      ↪ "spendhealth": world['spendhealth'],
      ↪ "fertility": world['fertility']}
pred_residual_corr = pd.DataFrame(df)
print(pred_residual_corr.corr())
```

	gdp_10_thou	pop_total	spendhealth	fertility
gdp_10_thou	1.000000	-0.036051	0.661569	-0.392485
pop_total	-0.036051	1.000000	-0.125908	-0.062368
spendhealth	0.661569	-0.125908	1.000000	-0.404251
fertility	-0.392485	-0.062368	-0.404251	1.000000

- k) Health spending and the GDP seem to be highly correlated. This could be a problem in interpreting the model in (b), because it means GDP could be a confounding variable between health spending and life expectancy. Therefore, we can't determine whether health spending has a causal effect on life expectancy.

6 Citations

- 1) <https://numpy.org/doc/stable/reference/routines.math.html>; used this source to find various numpy functions to help work on Q3
- 2) <https://numpy.org/doc/stable/reference/generated/numpy.percentile.html#numpy.percentile>; used this source to find the percentile method in numpy
- 3) <https://www.geeksforgeeks.org/python-pandas-dataframe-corr/>; used this source to learn how to use the corr() method in pandas for Q5.